

L2L Framework: Prompt Context Engineering

Design Rationale, Tier Architecture, and Production Decision Log

Kevin Linares

2026-03-23

Overview

This document describes the prompt construction process for the Latent-to-Language (L2L) item imputation framework. It records the design rationale, tier structure, prompt architecture, iterative development decisions, and production results for peer review. No API calls are made here — all prompts are printed from `item_prompt_functions.R` for inspection.

The L2L framework tests whether translating psychometric model *conclusions* — not raw parameters — into natural language prompts enables LLMs to impute masked survey item responses competitively with tree-based methods. The production design uses a **model-then-adjust** paradigm: IRT/CFA model predictions serve as starting points that LLMs refine based on residual patterns, correlated-item signals, and demographic probability context.

Key design principle: Prompts communicate model conclusions in natural language. The LLM is a reasoner, not a calculator. It never sees discrimination parameters, factor loadings, Q3 matrices, or threshold tables.

All prompt construction functions are in `item_prompt_functions.R`, a self-contained R script sourced by Script 04:

```
source("D:/repos/UMD_classes_code/TSE_II_SURV721/code/item_prompt_functions.R")
```

Tier Structure

Output Format

Batch prompting with model-specific batch sizes: Claude Sonnet = 50, Gemma 27B = 25, Maverick = 25. Each respondent block shows demographics, observed items with verbal labels and [MISSING] markers, and tier-specific psychometric context. Output is one line per respondent: `row_id: item1=value, item2=value`.

LLM Tiers

- **S (Demographics only):** Demographics + observed trust items as verbal labels + [MISSING] markers. Pure world-knowledge prediction. No model predictions, no psychometric context.
- **P1 (CFA augmented):** S + per-respondent CFA factor score (percentile and verbal label) + per-masked-item CFA prediction anchor + 2 most correlated observed items with CFA residuals + overall residual pattern summary.
- **P2 (IRT augmented):** S + per-masked-item IRT modal prediction with verbal confidence level and top-2 category probabilities + 2 most correlated observed items with IRT residuals + overall IRT residual pattern + demographic probability context (P(high trust | group) vs. population).

Traditional Method Tiers (Script 03)

- **S:** Demographics + observed trust items (masked as NA)
- **P1:** S + CFA factor score (from partial data)
- **P2:** S + IRT θ (from partial data, no SE)

Demographics (Screened by Script 02)

Demographics were screened using eta-squared with both IRT theta and CFA factor score (threshold: $\eta^2 \geq 0.005$). Each survey retains only demographics that pass screening:

- **LAPOP Mexico:** age_cat, Edu, Urban, Employment (male dropped, $\eta^2 = 0.0005$ with θ , 0.0007 with CFA)
- **LB Colombia:** age_cat, Edu, Employment (male and Urban dropped)

Psychometrician Persona

The system prompt uses a specific persona: a survey psychometrician with a PhD from JPSM at the University of Maryland, specializing in survey data science. The persona establishes domain expertise in institutional trust measurement for Latin American populations.

The system prompt frames the task as **refinement**, not binary keep/adjust:

“Your job is to **REFINE** these predictions using information the model cannot capture. The model only sees numerical patterns — you can see response consistency, demographic context, and institutional meaning.”

Reasoning Rules (System Prompt)

1. **Residual pattern is the strongest signal.** If the respondent consistently rates above or below the model on observed items, adjust in that direction.
2. **Correlated items provide direct evidence.** If the two most related observed items are both above (or below) the model, adjust in the same direction.
3. **Demographic context is a sanity check.** Respondent-level patterns always outweigh group averages.
4. **Domain knowledge matters.** Trust in armed forces is substantively different from trust in the judiciary.
5. **Low confidence = your reasoning matters most.** When the model’s top category has < 30% probability, do not simply repeat the prediction.
6. **Halo effects.** If a respondent rates all observed institutions consistently high or low, missing items almost certainly follow.
7. **Integer output.** Predictions must be on the original scale.

The Model-Then-Adjust Paradigm

Evolution from Raw Parameters

Early designs (V1) provided raw psychometric parameters: CFA standardized loadings, IRT discrimination and threshold tables in mirt format, Q3 residual correlation matrices, DIF flags from lordif, measurement invariance results, and theta regression coefficients. LLMs received no model predictions — they were asked to compute predictions from these parameters.

Result: Consistently below IRT baseline. LLMs cannot mentally evaluate the GRM probability function or convert discrimination parameters into item-level predictions.

Current Design (V2)

The model-then-adjust paradigm provides the statistical model’s conclusion as a starting point. For each masked item, the LLM receives:

1. **Anchor prediction:** The IRT modal category (P2) or CFA rounded prediction (P1).
2. **Confidence signal (P2 only):** Verbal label (HIGH/MODERATE/LOW) with top-2 category probabilities. Directive framing:
 - HIGH: “model is fairly certain”
 - MODERATE: “use residual pattern to refine”
 - LOW: “model is uncertain, your reasoning matters most”

3. **Correlated items (2 per masked item):** The two observed items most correlated with the masked item (from CFA model-implied correlation matrix), showing observed value, model expected value, and residual with explicit directional implication:
 - “+0.8 above → suggests this item should also be ABOVE model”
 - “-0.5 below → suggests this item should also be BELOW model”
4. **Overall residual pattern:** One-sentence prescriptive summary:
 - “This respondent consistently rates ABOVE the model (+0.5 avg) → adjust predictions UPWARD”
5. **Demographic probability context (P2 only):** Per-item model-based probability: “For your group, P(high trust on armed_forces) = 48% vs. population 43%.”

What Was Dropped from V1

- IRT parameter tables (discrimination a , thresholds b)
- CFA standardized loadings and SE
- Q3 local dependency matrix
- Conditional expectation $E[X_M|X_O]$
- DIF flags from lordif
- Measurement invariance non-invariant items
- Theta-demographic regression coefficients and R^2
- Theta with SE and percentile as standalone prompt line
- Chain-of-thought reasoning tags

All dropped elements either added noise (LLMs cannot compute from raw parameters) or were replaced by more interpretable alternatives (DIF flags → demographic probability context; theta regression → group-level P(high trust)).

Few-Shot Examples

Three training-fold exemplars per tier (P1, P2), spanning the latent trust spectrum. No few-shot for S tier (pure world knowledge).

1. **Low trust, low residual → KEEP.** Model confidence is high and the respondent’s pattern is consistent with predictions.
2. **Moderate trust, strong positive residual → ADJUST UP.** The respondent consistently rates above the model, and uncertainty is high.
3. **High trust, negative residual → ADJUST DOWN.** Despite high overall trust, this item’s pattern suggests lower.

Each example shows the complete prompt block for one respondent (demographics, observed items, residual summary, model prediction with confidence, correlated items) and the correct output line. Exemplars are selected from the training fold: low from the bottom tertile of latent scores, moderate from the middle, high from the top.

Reduction from 5 to 3 examples: The original 5-exemplar design included “moderate residual — small adjustment” and “extreme latent score — edge case.” Testing showed these two added tokens without improving accuracy. The three remaining examples cover the decision space: keep, adjust up, adjust down.

Multiple Imputation Design ($m = 5$)

Rationale

LLM predictions at P2 are treated as multiple imputations, mirroring the $m = 5$ majority-vote design used for mixgb and miceRanger in Script 03. Each fold $\times$ tier $\times$ model combination is run 5 times at graduated temperatures to introduce controlled stochasticity:

Run	Temperature
1	0.01
2	0.10
3	0.20
4	0.30
5	0.40

Benefits

1. **Parse failure recovery:** If a respondent fails to parse in 1 of 5 runs, the other 4 provide predictions.
2. **Prediction stabilization:** Majority vote across 5 runs reduces stochastic variance.
3. **Methodological symmetry:** Both traditional and LLM methods use $m = 5$ with majority vote.
4. **Between-imputation variance:** Raw per-run results are saved for MI variance analysis (RQ1 latent variable recovery).

Cost

Approximately $5\times$ the single-run cost. Claude Sonnet at batch 50: $\sim$ \$37 for all 5 folds $\times$ 3 tiers $\times$ 5 runs. Gemma and Maverick through OpenRouter are substantially cheaper.

Measurement Models

CFA (for P1)

Per fold, the CFA provides factor scores from partial data via `lavPredict()` with FIML for missing items. Factor scores are expressed as percentiles and verbal labels (VERY LOW through VERY HIGH) relative to the training distribution.

CFA predictions for masked items use the linear model: $\hat{X}_j = \hat{\tau}_j + \hat{\lambda}_j \hat{\eta}$, clamped to the scale range and rounded.

CFA residuals (observed – expected) for observed items provide the respondent-level consistency signal.

GRM/IRT (for P2)

The Graded Response Model provides θ from partial data via `fscores()` (EAP method). For each masked item, `probtrace()` computes the full category probability vector, from which the modal category, confidence level, and top-2 probabilities are derived.

IRT residuals (observed – $E[X_j|\hat{\theta}]$) for observed items are the primary adjustment signal.

For LB (descending scale), IRT reverse-coding ensures higher θ = higher trust. Probabilities are reversed back to the original scale in prompts.

Demographic Profiling (Script 02)

Script 02 computes model-based probabilities per demographic group: $P(X_j \geq k|\bar{\theta}_{\text{group}})$ using the IRT model. These replace the DIF flags and theta regression from V1 with directly actionable probability comparisons. In the prompt: “For your group, P(high trust on armed_forces) = 48% vs. population 43%.”

Baselines

- **IRT baseline:** $\arg \max_k P(X_j = k|\hat{\theta})$ from partial-data theta
- **CFA baseline:** $\text{round}(\hat{\tau}_j + \hat{\lambda}_j \hat{\eta})$ clamped to scale

Both are computed from partial data and recorded alongside every prediction in the results tibble.

Setup

```

library(lavaan)
library(mirt)
library(tidyverse)

set.seed(721)

code_dir    <- "D:/repos/UMD_classes_code/TSE_II_SURV721/code"
results_dir <- "D:/repos/UMD_classes_code/TSE_II_SURV721/results"

# Source item imputation prompt functions
source(file.path(code_dir, "item_prompt_functions.R"))

# Load fold data and psychometric objects
load(file.path(results_dir, "fold_data.RData"))
load(file.path(results_dir, "cfa_reporting.RData"))
load(file.path(results_dir, "irt_models.RData"))

cat("Loaded: folds_list, cfa_list, irt_list, masking_tbl,\n")

```

Loaded: folds_list, cfa_list, irt_list, masking_tbl,

```
cat("  cfa_params_tbl, irt_params_tbl, q3_pairs_tbl,\n")
```

cfa_params_tbl, irt_params_tbl, q3_pairs_tbl,

```
cat("  demo_profile_tbl, irt_models\n")
```

demo_profile_tbl, irt_models

```

# Helper: wrap long lines for PDF output
wrap_prompt <- function(text, width = 78) {
  lines <- str_split(text, "\n")[[1]]
  map_chr(lines, function(line) {
    if (nchar(line) <= width) return(line)
    str_wrap(line, width = width, exdent = 2)
  }) |> str_c(collapse = "\n")
}

```

Survey Configuration

```

config <- list(
  lapop = list(
    survey_label      = "a nationally representative survey conducted in Mexico in
↪ 2023",
    demo_vars         = c("age_cat", "Edu", "Urban", "Employment"),
    trust_items       = c(
      "justice_system", "electoral_tribunal", "armed_forces",
      "legislature", "public_ministry", "police", "auditor",
      "political_parties", "supreme_court", "municipality",
      "media", "elections"
    ),
    factor_names      = "Trust",
    trust_ascending   = TRUE,
    item_scale_desc   = "1 = no trust at all, 7 = a lot of trust",
    cfa_syntax = '
      Trust =~ justice_system + electoral_tribunal + armed_forces +
               legislature + public_ministry + police + auditor +
               political_parties + supreme_court + municipality +
               media + elections
      armed_forces ~~ legislature
      public_ministry ~~ auditor
    ,
  ),
  lb = list(
    survey_label      = "a nationally representative survey conducted in Colombia
↪ in 2023",
    demo_vars         = c("age_cat", "Edu", "Employment"),
    trust_items       = c(
      "Congress", "Natl_Govt", "Judiciary", "Parties",
      "Electoral_Inst", "President"
    ),
    factor_names      = "Political",
    trust_ascending   = FALSE,
    item_scale_desc   = "1 = a lot of trust, 4 = no trust at all",
    cfa_syntax = '
      Political =~ Congress + Natl_Govt + Judiciary + Parties +
                  Electoral_Inst + President
      Natl_Govt ~~ President
    ,
  )
)

cat("LAPOP:", length(config$lapop$trust_items), "items,",
    config$lapop$item_scale_desc, "\n")

```

LAPOP: 12 items, 1 = no trust at all, 7 = a lot of trust


```
cat("LB:   ", length(config$lb$trust_items), "items,",  
    config$lb$item_scale_desc, "\n")
```

LB: 6 items, 1 = a lot of trust, 4 = no trust at all

System Prompt

The system prompt is tier-invariant — the same persona and reasoning rules apply across all tiers. Only the user prompt content changes.

```
system_prompt <- build_system_prompt(config$lapop)  
cat(wrap_prompt(system_prompt))
```

You are a survey psychometrician with a PhD from the Joint Program in Survey Methodology (JPSM) at the University of Maryland, specializing in survey data science. You are reviewing an institutional trust scale from a nationally representative survey conducted in Mexico in 2023 where some item responses are missing and need to be imputed.

Scale: 1 = no trust at all, 7 = a lot of trust

Your task: For each respondent, a statistical model provides a starting-point prediction for each missing item. Your job is to REFINER these predictions using information the model cannot capture. The model only sees numerical patterns - you can see response consistency, demographic context, and institutional meaning.

HOW TO REASON:

1. Start from the model prediction, then look at the RESIDUAL PATTERN. If this respondent consistently rates ABOVE or BELOW the model on observed items, the missing items almost certainly follow the same direction. Adjust accordingly - this is the strongest signal.
2. Check the CORRELATED ITEMS. If the two most related observed items are both above (or below) the model, that is direct evidence to adjust in the same direction.
3. Consider DEMOGRAPHIC CONTEXT as a sanity check. If the model predicts a value that is unusual for this respondent's demographic group, that may warrant a small adjustment - but respondent-level patterns always outweigh group averages.
4. Use DOMAIN KNOWLEDGE. Trust in armed forces is substantively different from trust in the judiciary. Consider what you know about how these specific institutions are perceived in this country.
5. When model confidence is LOW (the top category has <30% probability), the model is genuinely uncertain - your contextual reasoning adds the most value

- here. Do not simply repeat the model prediction.
6. Trust scales exhibit HALO EFFECTS. If a respondent rates all observed institutions consistently high or low, the missing items almost certainly follow suit.
 7. Your prediction must be an INTEGER on the scale (1 = no trust at all, 7 = a lot of trust).

OUTPUT: One line per respondent. Format:
row_id: item1=value, item2=value
No explanation. Only prediction lines.

Prompt Examples

Tier S (No Model)

```
# Pick one test respondent from LAPOP fold 1
test_row <- folds_list$lapop |>
  filter(fold_id == 1, split == "test") |>
  slice(1)

masked <- c("armed_forces", "legislature")

info_s <- list(
  rid          = test_row$row_id,
  masked_items = masked
)

s_block <- build_respondent_block_s(info_s, test_row, config$lapop)
cat(wrap_prompt(s_block))
```

RESPONDENT 18:

Demographics: 30-44 | Secondary | Urban | Employed

Observed trust ratings:

justice_system = 5 (Moderate trust)
electoral_tribunal = 5 (Moderate trust)
armed_forces = [MISSING]
legislature = [MISSING]
public_ministry = 5 (Moderate trust)
police = 2 (Very little trust)
auditor = 5 (Moderate trust)
political_parties = 4 (Some trust)
supreme_court = 4 (Some trust)
municipality = 5 (Moderate trust)
media = 5 (Moderate trust)
elections = 4 (Some trust)

Predict:

armed_forces: [MISSING]
legislature: [MISSING]

Key Design Decisions

Model-Then-Adjust vs. Predict-From-Scratch

LLMs receiving the IRT prediction as an anchor consistently outperformed LLMs receiving only model parameters. LLMs cannot mentally evaluate the GRM probability function, but they can reason about whether to adjust a given prediction based on contextual signals. This is the core insight of the L2L framework: **translate model conclusions, not model parameters**.

Conclusions Over Parameters

The V1 design included IRT discrimination tables (a_j, b_{jk}), CFA loadings (λ_j), Q3 matrices, DIF flags, and regression coefficients. LLMs treated these as noise. The V2 design replaced all parameters with natural language conclusions:

V1 (Parameters)	V2 (Conclusions)
$a_j = 3.30, b_{jk}$ table	“Model prediction: 5, Confidence: HIGH (5 at 42%, 4 at 28%)”
Q3 matrix	“Most informative: police = 6 (+0.8 above → adjust up)”
DIF flag: “media shows DIF by Edu”	“For your group, P(high trust on media) = 35% vs. population 43%”
$\theta \sim \text{Edu}: b = -0.15$	(Absorbed into demographic probability context)
Residual: -0.32	“Respondent rates BELOW model → adjust DOWNWARD”

Prescriptive vs. Descriptive Prompts

V1 prompts were descriptive: “This respondent is +0.5 above model expectations.” V2 prompts are prescriptive: “This respondent consistently rates ABOVE the model (+0.5 avg) → adjust predictions UPWARD.” Testing showed that prescriptive framing does not significantly improve accuracy over descriptive framing for Claude Sonnet, but it reduces the gap for smaller models (Gemma, Maverick) by making the adjustment direction explicit.

No Chain-of-Thought

CoT was tested with both Maverick and Claude Sonnet under both V1 and V2 architectures. In every configuration, CoT degraded accuracy. Parse rates dropped and predictions were worse. The reasoning step introduces noise — the LLM talks itself out of correct adjustments.

No Base Rates

Base rate distributions encourage modal parroting. The IRT model predictions implicitly encode population-level information. Explicit base rates are redundant and harmful.

Three Few-Shot Examples (Reduced from Five)

The five-exemplar design was reduced to three (low/moderate/high trust) without performance loss. The two dropped exemplars (moderate residual and extreme latent score) added tokens without covering distinct adjustment scenarios.

Batch Size

Testing across sizes (1, 5, 10, 20, 25, 50) showed negligible accuracy differences for Claude Sonnet (range: 62.3–63.9% on LB P2). Claude uses batch 50 for cost efficiency. Gemma and Maverick use batch 25 due to higher parse failure rates at larger sizes.

Masking Design

1–4 items per respondent, randomly assigned (`sample(1:4)` count, `sample(items, size = k)` selection). Produces a continuous degradation curve. Masking table generated once in Script 02 and shared across all methods.

Opus 4.6 vs. Sonnet 4.6

Opus was tested on one fold — accuracy matched Sonnet at approximately 15× the cost. Sonnet is the production model.

Models Evaluated

LLMs

- **Claude Sonnet 4.6** (Anthropic API, batch 50): Strongest LLM. The only model consistently above IRT baseline at P2 on both surveys.
- **Llama 4 Maverick** (Meta, OpenRouter, batch 25): Below IRT at P2 but shows the tier ordering ($S < P2$).
- **Gemma 3 27B** (Google, OpenRouter, batch 25): Weakest LLM. Below IRT at P2 on both surveys.

Traditional Benchmarks

- **mixgb** (XGBoost, $m = 5$ majority vote): Strongest traditional benchmark.
- **miceRanger** (Random forest, $m = 5$ majority vote): Below mixgb on both surveys.

Parse Rates (5-Fold Production, $m=5$ MI)

Method	LAPOP S	LAPOP P1	LAPOP P2	LB S	LB P1	LB P2
Claude	90.1%	100%	100%	86.1%	100%	100%
Gemma	99.9%	96.9%	98.1%	100%	96.1%	98.6%
Maverick	100%	100%	100%	100%	100%	100%

Claude’s lower S-tier parse rate reflects the absence of the model anchor — without a structured prediction block, Claude occasionally produces verbose output. P1 and P2 parse rates are 100% for Claude.

Production Results (5-Fold CV, $m=5$ MI)

Accuracy by Method $\times$ Tier $\times$ Survey

Method	LAPOP S	LAPOP P1	LAPOP P2	LB S	LB P1	LB P2
Claude	.362	.379	.394	.603	.603	.614
Maverick	.366	.371	.367	.595	.581	.592
Gemma	.375	.347	.358	.589	.557	.586
IRT	—	—	.388	—	—	.600
mixgb	.382	.393	.393	.566	.574	.589

Claude P2 exceeds IRT by +0.6pp (LAPOP) and +1.4pp (LB). Claude P2 exceeds mixgb P2 by +0.1pp (LAPOP) and +2.5pp (LB). The advantage is scale-dependent: larger on the 4-point LB scale.

Note: Traditional method results at P1/P2 are reported for completeness but the fair comparison is LLM P1/P2 vs. traditional S, since trees at P1/P2 see the complete training response matrix plus psychometric features — an information advantage that has nothing to do with the L2L innovation. See Script 05 (H1c) for the fair comparison.

Key Findings

1. **Tier ordering (H1a):** Claude shows monotonic $S < P1 < P2$ on both surveys. Gemma and Maverick show the ordering on some but not all conditions.
2. **Scale dependency:** LLM advantage is larger on the 4-point scale (LB) where categorical predictions have higher base rates. On the 7-point scale (LAPOP), the IRT baseline is harder to beat.
3. **Model capability threshold:** Only the frontier model (Claude Sonnet) consistently exceeds IRT. Gemma 27B and Maverick cannot reliably extract the adjustment signal from residual patterns. This is itself a finding about the reasoning capability required for psychometric context engineering.
4. **Parse rate $\rightarrow$ data quality:** MI replication ($m = 5$) recovers from parse failures. After majority vote, effective coverage is $> 98\%$ for all model $\times$ tier conditions.

Prompt Engineering Evolution

1. **V1 (Raw parameters):** CFA loadings, IRT discrimination/ thresholds, Q3, DIF flags, MI results, theta regression, few-shot reasoning chains. No model predictions as anchors. Result: consistently below IRT baseline.
2. **V2a (Model-then-adjust, passive):** IRT/CFA predictions as anchors with “default to keeping unless clear evidence.” Five few-shot examples. DIF cues and conditional expectations included. Result: LLMs anchored too hard — zero marginal recovery.
3. **V2b (Model-then-adjust, simplified):** Dropped all raw parameters, DIF, MI, conditional expectations. Added demographic probability context from Script 02. Reduced to 3 few-shot examples. Result: Claude matches IRT on LAPOP, exceeds on LB.
4. **V2c (Prescriptive framing):** Changed residual summaries and correlated items from descriptive to prescriptive (“ $\rightarrow$ adjust UPWARD”). Changed system prompt from “decide whether to KEEP or ADJUST” to “your job is to REFINE.” Dropped the ± 2 cap. Result: no significant improvement for Claude, marginal improvement for smaller models. Adopted as production.
5. **MI replication ($m = 5$):** Graduated temperature schedule (0.01–0.40) across 5 runs with majority vote. Matches traditional MI design. Recovers parse failures and stabilizes predictions.

Limitations

1. **Listwise deletion** reduces samples (LAPOP: 1,430; LB: 1,073).
2. **ML estimation** treats ordinal items as continuous.
3. **Survey weights not used** — evaluation is case-level.
4. **Unidimensional IRT** — Q3 identifies local dependency but does not model it.
5. **7-point LAPOP** produces diffuse IRT distributions ($\max P < 0.25$), limiting the confidence signal.
6. **CFA refit per fold** in Script 04 introduces minor variation from Script 02 CFA.
7. **Model-then-adjust** bounds accuracy near IRT — marginal recovery is modest (+0.6pp LAPOP, +1.4pp LB for Claude).
8. **Random masking (1–4)** introduces heterogeneity in latent score estimation quality.
9. **Cost constraints** — 3 LLMs $\times$ 3 tiers $\times$ 2 surveys $\times$ 5 folds $\times$ 5 MI runs.
10. **Batch size asymmetry** — Claude uses batch 50, others use 25. Sensitivity analysis shows negligible effect of batch size on accuracy.
11. **S-tier parse rate** — Claude’s lower S-tier parse rate (86–90%) is a limitation that could bias S-tier results toward parseable (potentially easier) cases.

Verify item_prompt_functions.R

```
source_path <- file.path(code_dir, "item_prompt_functions.R")

if (file.exists(source_path)) {
  tryCatch({
    source(source_path, local = TRUE)
    cat("Verified item_prompt_functions.R at:", source_path, "\n")
    cat("File size:", round(file.size(source_path) / 1e3, 1), "KB\n")
    cat("Functions loaded successfully.\n")
  }, error = function(e) {
    cat("ERROR sourcing:", e$message, "\n")
  })
} else {
  cat("WARNING: not found at:", source_path, "\n")
}
```

Verified item_prompt_functions.R at: D:/repos/UMD_classes_code/TSE_II_SURV721/code/item_prompt.
File size: 30.7 KB
Functions loaded successfully.